

# Assignment 3 Report

Aditya Jain

## 1 Part A

### 1.1 Part A.1

An Intrusion Detection System (IDS) checks, monitors, and analyzes unauthorized and malicious attempts to gain access to a system or a network. Yes, we can implement IDS on the KDD Cup 99 dataset. The workflow for implementing an IDS is as follows:

1. Data processing: The first step is to process and standardize the input data. The categorical data is converted to numerical entries, followed by scaling to zero mean and unit variance.
2. Feature selection: Chi-Squared test of independence is performed to remove redundant and irrelevant features from the dataset. This step gives us important features that should be used for model training.
3. Model training: Support vector machine (SVM) is used for classification, which classifies the data into different classes by an N-dimensional hyperplane.
4. Model evaluation: Finally, the model's accuracy is evaluated on the test dataset.

### 1.2 Part A.2

Since retrieving from *urllib* was throwing an error, I directly uploaded the data to Databricks filesystem (as explained in the lecture). Figure 1 shows the display result.

The screenshot shows the Databricks command history and a file display. The first command is `import urllib.request`. The second command is `display(dbutils.fs.ls("dbfs:/FileStore/tables/kddcup_data_10_percent.gz"))`. Below the command, a table shows the file details:

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/kddcup_data_10_percent.gz	kddcup_data_10_percent.gz	2144903	1677090361000

Below the table, it says "1 row | 0.70 seconds runtime" and "Refreshed now".

Figure 1: Result of Part A.2

### 1.3 Part A.3

Figure 2 shows the 10 values of the data and its type.

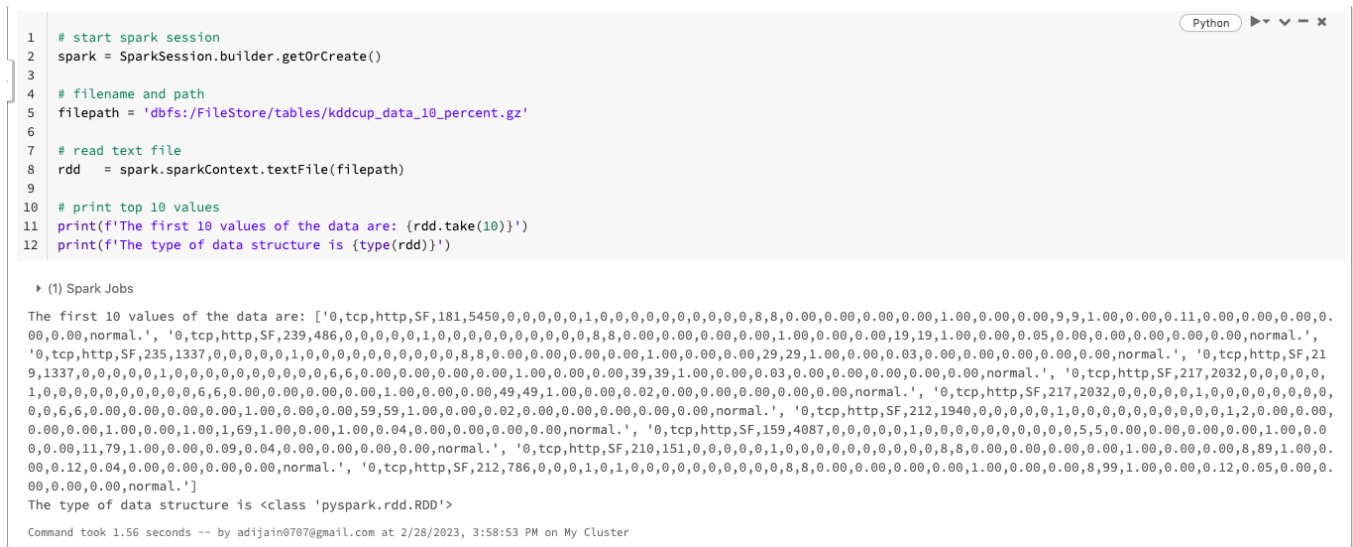


Figure 2: Result of Part A.3

## 1.4 Part A.4

Figure 3 shows an example entry of the data. There are a total of 41 features, with an additional column determining the label of connection type (normal or attack). There are a total of 494,021 row entries in the data.

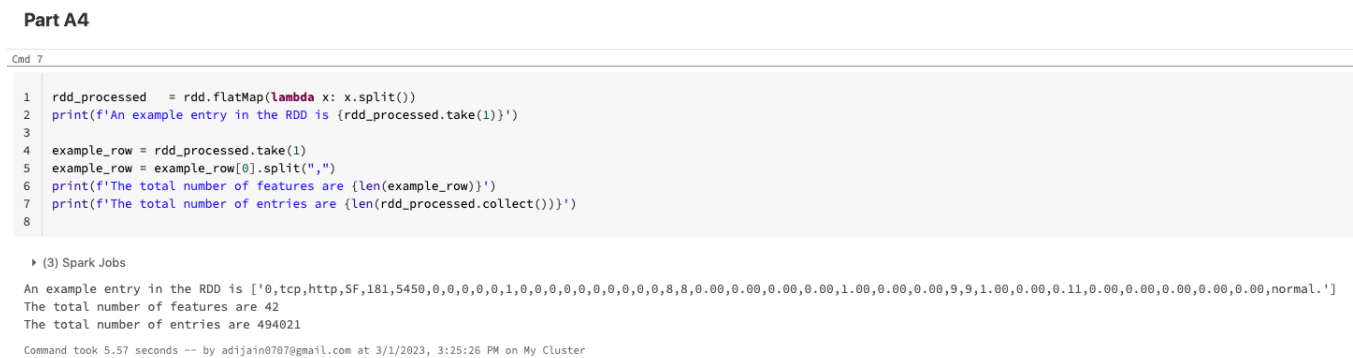


Figure 3: Result of Part A.4

## 1.5 Part A.5

A new RDD and dataframe is built using the given six columns and label. Figure 4 shows the dataframe schema and ten values.

```

1 rdd_six_columns = rdd_processed.map(lambda p: Row(
2                                     duration = int(p.split(",")[0]),
3                                     protocol_type = str(p.split(",")[1]),
4                                     service = str(p.split(",")[2]),
5                                     flag = str(p.split(",")[3]),
6                                     src_bytes = int(p.split(",")[4]),
7                                     dst_bytes = int(p.split(",")[5]),
8                                     label = str(p.split(",")[-1])
9                                     ))
10
11 df = rdd_six_columns.toDF()
12 df.printSchema()
13 df.show(10)

```

▶ (2) Spark Jobs

df: pyspark.sql.dataframe.DataFrame = [duration: long, protocol\_type: string ... 5 more fields]

```

root
|-- duration: long (nullable = true)
|-- protocol_type: string (nullable = true)
|-- service: string (nullable = true)
|-- flag: string (nullable = true)
|-- src_bytes: long (nullable = true)
|-- dst_bytes: long (nullable = true)
|-- label: string (nullable = true)

```

duration	protocol_type	service	flag	src_bytes	dst_bytes	label
0	tcp	http	SF	181	5450	normal.
0	tcp	http	SF	239	486	normal.
0	tcp	http	SF	235	1337	normal.
0	tcp	http	SF	219	1337	normal.
0	tcp	http	SF	217	2032	normal.
0	tcp	http	SF	217	2032	normal.
0	tcp	http	SF	212	1940	normal.
0	tcp	http	SF	159	4087	normal.
0	tcp	http	SF	210	151	normal.

Figure 4: Result of Part A.5

## 1.6 Part A.6

Figure 5 shows the total number of connections based on *protocol\_type*.

Figure 6 shows the total number of connections based on *service* and figure 7 is the corresponding bar graph.

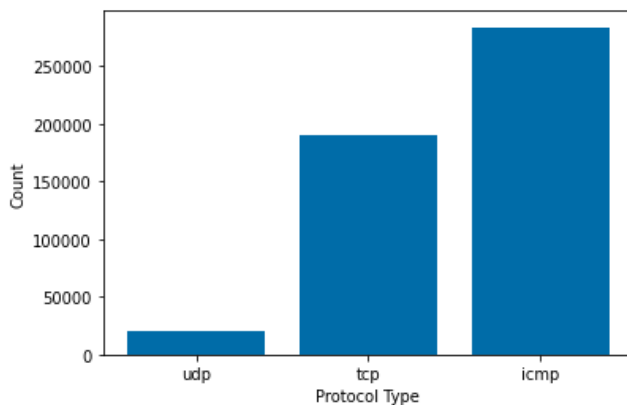
```

1  # analysis based on protocol type
2  data_protocol_type = df.rdd.map(lambda x: x.protocol_type).collect()
3
4  dictionary = {}
5  for item in data_protocol_type:
6      if item not in dictionary.keys():
7          dictionary[item] = 1
8      else:
9          dictionary[item] += 1
10
11 dictionary = dict(sorted(dictionary.items(), key=lambda item: item[1]))
12 print(f'The total number of connections based on protocol_type are {dictionary}')
13
14 # plot
15 names = list(dictionary.keys())
16 values = list(dictionary.values())
17
18 plt.bar(range(len(dictionary)), values, tick_label=names)
19 plt.xlabel('Protocol Type')
20 plt.ylabel('Count')
21 plt.show()

```

► (1) Spark Jobs

The total number of connections based on protocol\_type are {'udp': 20354, 'tcp': 190065, 'icmp': 283602}



Command took 14.12 seconds -- by adijain0707@gmail.com at 3/2/2023, 10:17:02 AM on My Cluster

Figure 5: Result of Part A.6 *protocol\_type*

## 1.7 Part A.7

Figure 8 shows the number of connections based on the type of connection. The 'normal' connections are only a small fraction and the others represent the attack type.

Figure 9 shows a bar plot for the counts of various error status of the connection.

Figure 10 is a comparison between average *src\_bytes* and average *dst\_bytes* transferred. It is evident that a lot more data is transferred from source to destination than other way round.

```

1 # analysis based on service type
2 data_service_type = df.rdd.map(lambda x: x.service).collect()
3
4 dictionary = {}
5 for item in data_service_type:
6     if item not in dictionary.keys():
7         dictionary[item] = 1
8     else:
9         dictionary[item] += 1
10
11 dictionary = dict(sorted(dictionary.items(), key=lambda item: item[1]))
12 print(f'The total number of connections based on service are {dictionary}')
13
14 # plot
15 names = list(dictionary.keys())
16 values = list(dictionary.values())
17
18 plt.figure(figsize=(8,14))
19 plt.barh(names, values)
20 plt.xlabel('Count')
21 plt.ylabel('Service Type')
22 plt.show()

```

▶ (1) Spark Jobs

The total number of connections based on service are {'pm\_dump': 1, 'tftp\_u': 1, 'red\_i': 1, 'tim\_i': 7, 'X11': 11, 'urh\_i': 14, 'IRC': 43, 'Z39\_50': 92, 'netstat': 95, 'ctf': 97, 'name': 98, 'kshell': 98, 'http\_443': 99, 'exec': 99, 'netbios\_dgm': 99, 'pop\_2': 101, 'ldap': 101, 'link': 102, 'netbios\_ns': 102, 'daytime': 103, 'efs': 103, 'login': 104, 'hostnames': 104, 'ssh': 105, 'nnsdp': 105, 'supdup': 105, 'uucp': 106, 'klogin': 106, 'vmnet': 106, 'bgp': 106, 'mtp': 107, 'sunrpc': 107, 'netbios\_ssn': 107, 'nntp': 108, 'courier': 108, 'printer': 109, 'whois': 110, 'sql\_net': 110, 'rje': 111, 'shell': 112, 'echo': 112, 'systat': 115, 'iso\_tsap': 115, 'domain': 116, 'discard': 116, 'gopher': 117, 'imap4': 117, 'remote\_job': 120, 'csnet\_ns': 126, 'time': 157, 'pop\_3': 202, 'auth': 328, 'ntp\_u': 380, 'telnet': 513, 'urp\_i': 538, 'finger': 670, 'ftp': 798, 'eco\_i': 1642, 'ftp\_data': 4721, 'domain\_u': 5863, 'other': 7237, 'smtp': 9723, 'http': 64293, 'private': 110893, 'ecr\_i': 281400}

Figure 6: Result of Part A.6 *service*

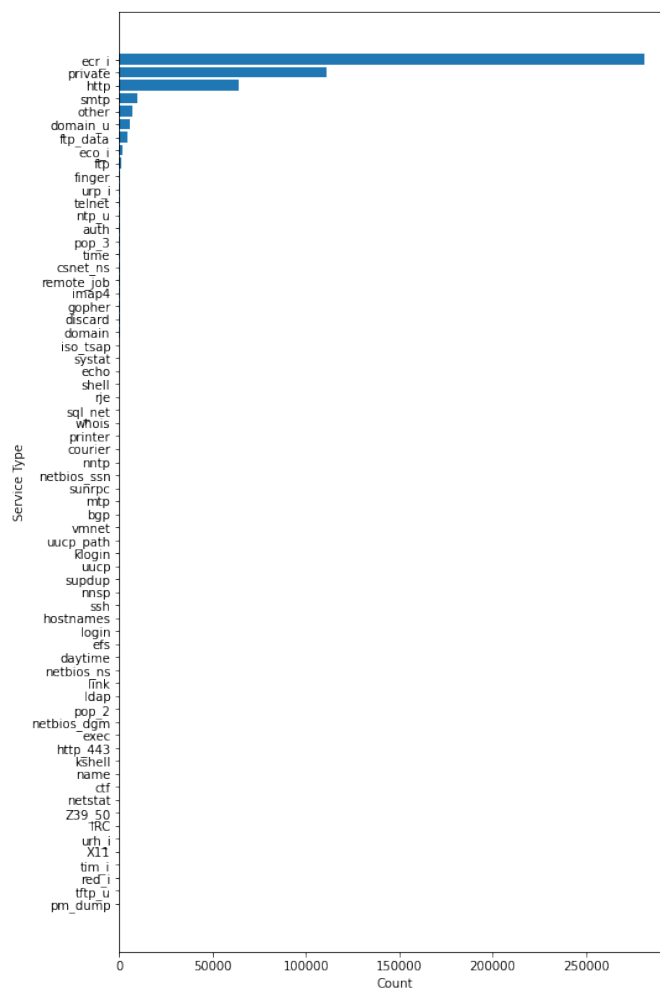


Figure 7: Result of Part A.6 *service* bar graph

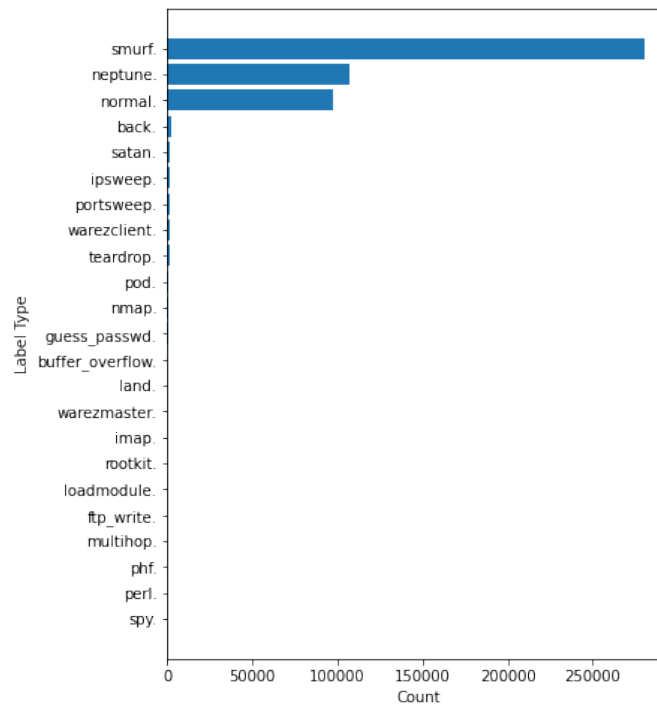


Figure 8: Counts of type of connection

## 1.8 Part A.8 and A.9

The six attributes from part 1.5 are used for training models. Figure 11 shows the data transformation that is applied before training. The categorical data is converted to numerical data and features are concatenated into a single vector. The data is split into training (80%) and test (20%) sets.

The two models trained are: **linear SVM classifier** (LinearSVC) and **multilayer perceptron classifier** (MLP). LinearSVC is a classical ML algorithm for classification, which constructs a hyperplane(s) in high-dimensional space between the datapoints of different classes. MLP is a deep-learning based algorithm that uses a neural network architecture for learning patterns in the data. The test classification accuracy for LinearSVC is 93.85% (figure 12) and 99% (figure 13) for MLP. The near perfect accuracy for MLP is attributed to its architecture: the structure of hidden layers make it theoretically an universal approximator. The other contributing factors are simplicity of the data and large number of epochs used for training.

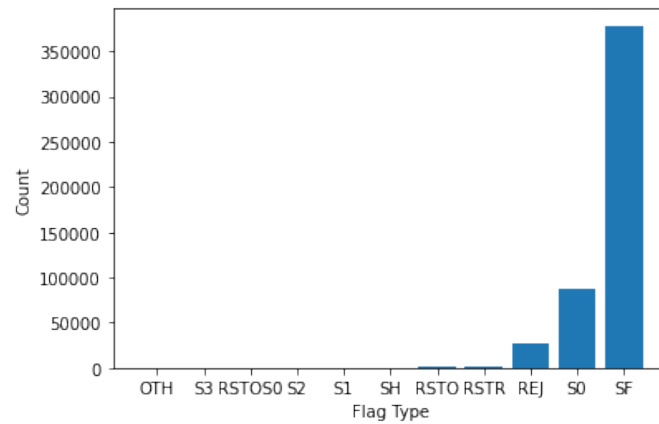


Figure 9: Comparison based on flag type

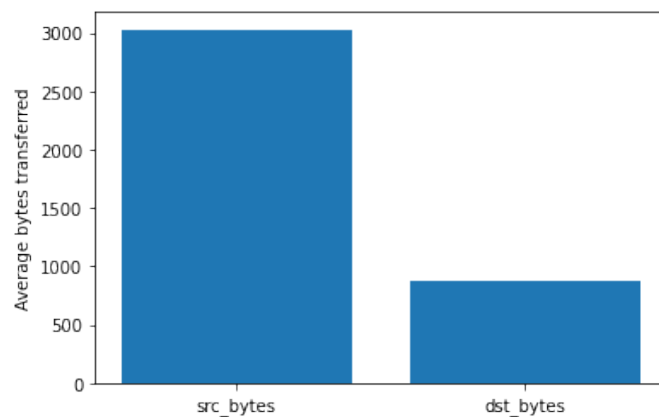


Figure 10: Data transfer

## 2 Part B

### 2.1 Part B.1

1. A platform as a service (PaaS) solution that hosts web apps in Azure provide professional development services to continuously add features to custom applications.

**Yes** - PaaS provides a framework where developers can build or customize cloud-based applications. Being able to add custom features to applications is one of the core features of PaaS.

2. A platform as a service (PaaS) database offering in Azure provides built-in high availability.

**Yes** - Just like other cloud computing services (SaaS, IaaS), high availability is one of the default components of PaaS.

### 2.2 Part B.2

A relational database must be used when: **Data will be stored as key/value pairs.**

A relational database is used when we have structured data and adheres to a strict schema. If the data needs to be stored in a structured format like a dictionary (key/value pairs), then relational database must be used.

**Data transformation**

Cmd 19

```

1 # club all attack types to one form
2 df_new = df.withColumn("label", when(df.label == "normal.", "normal").otherwise("attack"))
3
4 # convert categorical data to numeric data
5 categorical_columns = ['protocol_type', 'service', 'label', 'flag']
6 indexers = [StringIndexer(inputCol=column, outputCol=column+"_index").fit(df_new) for column in categorical_columns]
7 pipeline = Pipeline(stages=indexers)
8 df_new = pipeline.fit(df_new).transform(df_new)
9
10 # aggregate data into a feature vector
11 assembler = VectorAssembler(
12     inputCols=["duration", "protocol_type_index", "service_index", "flag_index", "src_bytes", "dst_bytes"],
13     outputCol="features")
14 df_ml_data = assembler.transform(df_new)
15 df_ml_data.show(5)
16
17 train_split = 0.8
18 test_split = 0.2
19 (training, test) = df_ml_data.randomSplit([train_split, test_split])

```

► (9) Spark Jobs

► df\_new: pyspark.sql.dataframe.DataFrame = [duration: long, protocol\_type: string ... 9 more fields]  
 ► df\_ml\_data: pyspark.sql.dataframe.DataFrame = [duration: long, protocol\_type: string ... 10 more fields]  
 ► training: pyspark.sql.dataframe.DataFrame = [duration: long, protocol\_type: string ... 10 more fields]  
 ► test: pyspark.sql.dataframe.DataFrame = [duration: long, protocol\_type: string ... 10 more fields]

duration	protocol_type	service	flag	src_bytes	dst_bytes	label	protocol_type_index	service_index	label_index	flag_index	features
0	tcp	http	SF	181	5450	normal	1.0	2.0	1.0	0.0	[0.0,1.0,2.0,0.0,...]
0	tcp	http	SF	239	486	normal	1.0	2.0	1.0	0.0	[0.0,1.0,2.0,0.0,...]
0	tcp	http	SF	235	1337	normal	1.0	2.0	1.0	0.0	[0.0,1.0,2.0,0.0,...]
0	tcp	http	SF	219	1337	normal	1.0	2.0	1.0	0.0	[0.0,1.0,2.0,0.0,...]
0	tcp	http	SF	217	2032	normal	1.0	2.0	1.0	0.0	[0.0,1.0,2.0,0.0,...]

only showing top 5 rows

Figure 11: Data transformation

**Model 1: Linear SVM**

Cmd 21

```

1 # import and fit model
2 lsvc = LinearSVC(labelCol="label_index", maxIter=20, regParam=0.1)
3 lsvcModel = lsvc.fit(training)
4
5 # print the coefficients and intercept for linear SVC
6 print("Coefficients: " + str(lsvcModel.coeficients))
7 print("Intercept: " + str(lsvcModel.intercept))
8
9 # evaluate on the test test
10 predictions = lsvcModel.transform(test)
11 evaluator = BinaryClassificationEvaluator(labelCol="label_index")
12 accuracy = evaluator.evaluate(predictions)
13 print(f'The classification accuracy of the Linear SVM model is {accuracy*100}')
```

► (87) Spark Jobs

► predictions: pyspark.sql.dataframe.DataFrame = [duration: long, protocol\_type: string ... 12 more fields]

Coefficients: [2.0849911498631425e-05,1.2201552395144577,6.753815757945768e-05,-1.2225838636634285,-4.495743192291753e-09,2.422045828482321e-06]

Intercept: -1.000648377787103

The classification accuracy of the Linear SVM model is 93.8490969752666

Command took 46.08 seconds -- by adjain0707@gmail.com at 3/4/2023, 4:46:41 PM on My Cluster

Figure 12: Linear SVM




**Model 2: Multilayer Perceptron Classifier**

Cmd 23

```

1  # number of layers in the model
2  layers = [6, 8, 7, 2]
3
4  # create the trainer and set its parameters
5  trainer = MultilayerPerceptronClassifier(maxIter=100, labelCol="label_index", layers=layers, blockSize=128)
6
7  # train the model
8  mlp_model = trainer.fit(training)
9
10 # evaluate on the test test
11 predictions = mlp_model.transform(test)
12 evaluator   = BinaryClassificationEvaluator(labelCol="label_index")
13 accuracy    = evaluator.evaluate(predictions)
14 print(f'The classification accuracy of the multilayer perceptron is {accuracy*100}')
```

► (94) Spark Jobs

►  predictions: pyspark.sql.dataframe.DataFrame = [duration: long, protocol\_type: string ... 13 more fields]

The classification accuracy of the multilayer perceptron is 99.00958167669113

Command took 4.74 minutes -- by adijain0707@gmail.com at 3/4/2023, 5:09:43 PM on My Cluster

Figure 13: Multilayer Perceptron Classifier

## 2.3 Part B.3

When you are implementing a Software as a Service solution, you are responsible for: **Defining scalability rules.**

In SaaS framework, the installation, configuration, and deployment of the software is taken by the service provider. The user needs to just define the scalability rules depending upon requirement.

## 2.4 Part B.4

1. To achieve a hybrid cloud model, a company must always migrate from a private cloud model.  
**No** - A hybrid cloud model is a mix of public cloud and on-premise deployment. Hence, to achieve a hybrid cloud model, a company needs to migrate only *some* of its infrastructure to a cloud service provider.

2. A company can extend the capacity of its internal network by using a public cloud.  
**Yes** - Through the hybrid cloud model, a company extend its on-premise capacity by using a public cloud.

3. In a public cloud model, only guest users at your company can access the resources in the cloud.  
**No** - Specific access privileges can be given to specific users of the company in the cloud.

## 2.5 Part B.5

A. A cloud service that remains available after a failure occurs - **Fault Tolerance**  
 Fault tolerance is the property of a cloud service in which it remains accessible after a failure.

B. A cloud service that can be recovered after a failure occurs - **Disaster Recovery**

Disaster recovery is the property of a cloud service in which data can be recovered after a failure.

C. A cloud service that performs quickly when demand increases - **Dynamic Scalability**

A cloud service has dynamic scalability if its infrastructure scales according to the demand.

D. A cloud service that can be accessed quickly from the internet - **Low Latency**

A cloud service has low latency if it can be accessed from the internet in a small amount of time.